



HØGSKOLEN
I BERGEN

BERGEN UNIVERSITY COLLEGE

Greedy algorithms and local search

Chapter 2





- Greedy algorithms
Step by step: Next step make some decision that is locally best possible
Typically primal infeasible
- Local search algorithms
Starts with an arbitrary feasible solution
Check if some small, local change results in improved objective function
- Both
 - popular and easy to implement
 - Good running time in practice



Scheduling jobs

- Scheduling jobs with deadlines on a **single** machine
 - n jobs, each job j
 - needs p_j time units of processing time
 - release date r_j
 - due date d_j
 - finish at time C_j
 - lateness $L_j = C_j - d_j$
 - one machine
 - can process at most one job at the time
 - must process a job to its completion
 - Minimize maximum lateness

$$L_{max} = \max_{j=1, \dots, n} L_j$$



Negative result

- Difficult to obtain near-optimal values
 - If there were a p -approximation algorithm, then for any input with optimal value 0, the algorithm must still find a schedule of objective function value at most $p \cdot 0 = 0$
 - Would imply $P = NP$
 - Consider the problem where all due dates are negative
 - Implies optimal value is always positive

Algorithm

- Earliest due date rule (EDD-rule)
 - Process next an available job with the earliest due date
- Theorem 2.2. The EDD-rule is a 2-approximation algorithm for the problem of minimizing the maximum lateness on a single machine subject to release dates with negative due dates.

Clustering

- Examples
 - Finding similarities and dissimilarities in large amount of data
 - Customers with similar purchasing behavior
 - Voting behavior
 - Search engines group webpages by similarity of topic



The k-center problem

- Input
 - Undirected, complete graph $G = (V, E)$
 - Distance $d_{ij} \geq 0$ between every pair of vertices $i, j \in V$. Smaller on previous slide means smaller d_{ij}
 - Integer k
- Assume
 - $d_{ii} = 0$ and $d_{ij} = d_{ji}$
 - Obey triangle inequality
- Goal: Choose k vertices (cluster centers) such that the maximum distance of a vertex to its cluster center is as small as possible



Notation

- Distance of a vertex i from a set $S \subseteq V$

$$d(i, S) = \min_{j \in S} d_{ij}$$



Greedy algorithm

Pick arbitrary $i \in V$

$S \leftarrow \{i\}$

while $|S| < k$ do

$j \leftarrow \arg \max_{j \in V} d(j, S)$

$S \leftarrow S \cup \{j\}$

-
- Theorem 2.3. The greedy algorithm is a 2-approximation algorithm for the k-center problem.

Negative result

- Theorem 2.4. There is no α -approximation algorithm for the k-center problem for $\alpha < 2$ unless $P = NP$.
- Consider the Dominating set problem which is NP-complete
 - Graph $G=(V,E)$ and integer k
 - Decide if there exists a set $S \subseteq V$ of size k such that each vertex is either in S , or adjacent to a vertex in S

Proof

- Given an instance of dominating set problem
- Define instance of k -center problem by setting the distance between adjacent vertices to 1 and nonadjacent vertices to 2
- Any p -approximation algorithm with $p < 2$ must always produce a solution of radius 1 if such a solution exists

Scheduling jobs

- n jobs to be processed
- m identical machines (running in parallel) to which each job may be assigned
- Each job $1, \dots, n$ must be processed on one of these machines for p_j time units without interruption
- Each job is available at time 0
- Each machine can process at most one job at the time
- The aim is to complete all jobs as soon as possible



Local search algorithm

- Start with any schedule
- loop
 - Consider the job ℓ that finishes last
 - if (job ℓ can be reassigned to another machine so it finishes earlier) {
 reassign it
 - else
 exit
- continue



Performance analysis

- Bounds of overall finish time
 - At least as long as the longest job

$$C_{max}^* \geq \max_{j=1, \dots, n} p_j$$

- At least one machine must have work at least corresponding to the average work per machine

$$C_{max}^* \geq \frac{1}{m} \sum_{j=1}^n p_j$$

Theorem

- Theorem 2.5. The local search procedure for scheduling jobs on identical parallel machines is a 2-approximation algorithm



Sketch of proof

- Let ℓ be the job that completes last in the schedule and let its completion time be C_ℓ
- Every other machine must be busy from time 0 until the start of job ℓ at time $S_\ell = C_\ell - p_\ell$
- Partition the schedule into two different time intervals 0 to S_ℓ and S_ℓ to C_ℓ .
- Both are less than C_{max}^* because of bounds on previous slide
- Have to show it is poly time (argue that we never transfer a job twice)



Greedy algorithm

- List scheduling algorithm
 - Assign the jobs as soon there is a machine available



-
- Theorem 2.6. The list scheduling algorithm for the problem of minimizing the makespan on m identical parallel machines is a 2-approximation algorithm.



-
- Longest processing time rule
 - Order the list with the longest jobs first



-
- Theorem 2.7. The longest processing time rule is a $4/3$ -approximation algorithm for scheduling jobs to minimize the makespan on identical machines.
 - Sketch of proof

Assume we have a counterexample
Show that such an example can not exist



Travelling Salesman problem (TSP)

- Given a set of cities $\{1, 2, \dots, n\}$
- Symmetric $n \times n$ cost matrix $C = (c_{ij})$ travelling from city i to city j .
- $c_{ii} = 0$ and $c_{ij} \geq 0$
- A feasible solution, or tour, traversal of the cities in the order $k(1), k(2), \dots, k(n)$ where $k(i)$ is the i 'th city to visit.
- Cost of tour: $c_{k(n)k(1)} + \sum_{i=1}^{n-1} c_{k(i)k(i+1)}$
- Want to find the tour with minimum cost



- NP-complete to decide whether a graph has a Hamiltonian cycle
- An approximation algorithm for TSP can be used to solve the Hamiltonian cycle in the following way. Set
 - $c_{ij} = 1$ if $(i, j) \in E$
 - $c_{ij} = \alpha n + 2$ otherwise
- Now an α -approximation algorithm could solve the Hamiltonian cycle problem



- Theorem 2.9. For any $\alpha > 1$, there does not exist an α -approximation algorithm for the travelling salesman problem on n cities, provided $P \neq NP$. In fact, the existence of an $O(2^n)$ -approximation algorithm for TSP would imply that $P = NP$.



Metric TSP

- Satisfy the triangle inequality

$$c_{ij} \leq c_{ik} + c_{kj}$$

- Will look at three approximation algorithms for metric TSP
 - Nearest addition algorithm
 - Double tree algorithm
 - Christofides' algorithm

Useful fact


- Lemma 2.10. For any input to the travelling salesman problem, the cost of the optimal tour is at least the cost of the minimum spanning tree on the same input
- Proof by contradiction. Assume it is not. Consider the tour. Remove the edge with larger weight. Then we have a spanning tree with cost less than a minimum spanning which is impossible.



Nearest addition algorithm

- Build a tour between the two closest cities, i and j .
- Tour $T = (i \rightarrow j \rightarrow i)$
- $S = \{i, j\}$
- while $(|S| < n)$
 - find a pair of cities $i \in S$ and $j \notin S$ for which cost c_{ij} is minimum
 - Let k be the city that follows i in current tour. Must consider the tour both clockwise and counterclockwise
 - Insert j in tour such that $T = (\dots \rightarrow i \rightarrow j \rightarrow k \rightarrow \dots)$

Figure





-
- Theorem 2.11. The nearest addition algorithm for the metric TSP is a 2-approximation algorithm
 - We notice that nodes are added in the same sequence as in Prim's algorithm for minimum spanning trees.



Proof of Theorem 2.11

- Idea. After each iteration the cost is at most twice the cost of the corresponding MST found by Prim
 - First iteration obvious
 - Later
 - Prim adds c_{ij}
 - We add $c_{ij} + c_{jk} - c_{ik}$
 - By triangle inequality $c_{jk} \leq c_{ji} + c_{ik}$ or $c_{jk} - c_{ik} \leq c_{ji}$
 - Hence increase is at most $c_{ij} + c_{ji} = 2c_{ij}$
 - Thus the cost of tour found is at most $2 \cdot \text{cost}(\text{Prim}) \leq 2 \cdot \text{OPT}$ and Theorem is proved



The double-tree algorithm

- (Saw it the first day)
- Find MST
- Double all the edges in the tree (then all nodes will have even degree)
- Find Eulerian tour (exist when all nodes have even degree. Easy to find)
- Consider tour. Skip nodes that are already visited. Because of triangle inequality cost can not increase

-
- Theorem 2.12. The double tree algorithm for the metric TSP is a 2-approximation algorithm
 - Proof: Same ideas as nearest addition algorithm

Perfect matching

- If we have $2k$ nodes, a perfect matching is k edges where all nodes appears exactly once
- It is possible to compute the perfect matching of minimum total cost in polynomial time.

Better algorithm

- If we sum all node degrees in a graph, the sum must be even (every edge contributes 2 in total degree)
- In double-tree algorithm, we doubled every edge in MSP so we could find an Eulerian tour
- However, the number of nodes of odd degree must be even
- Sufficient to find a minimum perfect matching among the nodes with odd degree and add these edges to the tree

Christofides' algorithm

- Find a MST
- Find a minimum perfect matching among the nodes with odd degree in MST. Add these edges to the MST
- Find a Eulerian tour in the resulting graph
- Consider the tour. Skip already visited vertices

Christofides' algorithm

- Theorem 2.13. Christofides' algorithm for the metric TSP is a $3/2$ -approximation algorithm



Proof of Christofides' algorithm

- Let O be set of odd degree vertices in MST
- Consider optimal tour of all vertices.
- Make a tour on just vertices in O where the vertices appears in the same order as in the optimal tour.
- This tour can not be longer than before because of the triangle inequality
- Color the edges red and blue, alternating colors as the tour is traversed
- The red and blue edges represent two perfect matchings. The cheapest must have cost at most $OPT/2$
- The minimum perfect matching can not be more expensive
- Hence Christofides' algorithm find a tour of cost at most $\text{cost(MST)} + \text{cost(Perfect matching)} \leq OPT + OPT/2 = 3/2OPT$

TSP, Negative results

- No better algorithm for the metric TSP is known
- Theorem 2.14. Unless $P = NP$, for any constant $\alpha < \frac{220}{219} \approx 1.0045$, no α -approximation algorithm for the metric TSP exists.
- Improved to $185/184 \approx 1.0054$



Maximizing Float in Bank Accounts

- Wish to open k bank accounts
- Let
 - B be the set of banks
 - P be the set of people we regularly pays
 - $v_{ij} \geq 0$ value created to be able to pay person $j \in P$ from bank account $i \in B$ (takes into account time, interest rate and other things)
- Wish to find
 - $S \subseteq B, |S| \leq k$ that maximizes $v(S) = \sum_{j \in P} \max_{i \in S} v_{ij}$



Greedy algorithm

- $S \leftarrow \emptyset$
- while $|S| < k$ do
 - $i \leftarrow \arg \max_{i \in B} v(S \cup \{i\}) - v(S)$
 - $S \leftarrow S \cup \{i\}$
- return S

Performance

- Theorem 2.16. The greedy algorithm gives a $(1 - \frac{1}{e})$ -approximation algorithm for the float maximization problem. ($e \approx 2.72$)
- Notice. Since this is a maximization problem, we have performance guarantee less than 1.



Finding minimum degree spanning trees

- Given a graph $G = (V, E)$
- Want to find a spanning tree T of G so as to minimize the maximum degree of nodes in T .



Finding minimum degree spanning trees

- Theorem 2.18. It is NP-complete to decide whether or not a given graph has a minimum-degree spanning tree of maximum degree two.
- Proof: We know that Hamiltonian path is NP-complete. A Hamiltonian path is a MSP where the maximum degree is two



Algorithm

- Start with an arbitrary spanning tree T
- Pick a vertex u .
- Look at all edges (v, w) that are not in T but if added to T creates a cycle containing u .
- Suppose $\max(d_T(v), d_T(w)) \leq d_T(u) - 2$
- Add (v, w) and remove edge (u, y) where (u, y) is part of the cycle (figure next slide)

50 Greedy algorithms and local search




Figure 2.5: Illustration of a local move for minimizing the maximum degree of a spanning tree. The bold solid lines are in the tree, and the dashed lines are graph edges not in the tree.



- Remember. For an algorithm to be an approximation algorithm
 - Must have a performance guarantee
 - Run in polynomial time
- Problem: to show that algorithm takes poly time
- Solution: apply local moves only to nodes whose degree is relatively high.
 - Let $\Delta(T)$ be the maximum degree of a node in T
 - Algorithm picks node that has degree at least $\Delta(T) - \ell$

-
- Theorem 2.19. Let T be a locally optimal tree. Then $\Delta(T) \leq 2OPT + \ell$, where $\ell = \lceil \log_2 n \rceil$



Polynomial time

- Idea
 - Defines a potential function
 - Defines a minimum potential
 - The potential drops by a certain percentage in each iteration
 - Can limit the number of iterations



Edge Coloring

- An undirected graph is k -edge-colorable if each edge can be assigned exactly one of k colors in such a way that no two edges with same colors share an endpoint
- Want to obtain a k -edge-coloring with k as small as possible



- Let Δ be the maximum degree of a vertex in the given graph
- Theorem 2.22. For graphs with $\Delta = 3$, it is NP-complete to decide whether the graph is 3-edge-colorable or not.



Algorithm that uses $\Delta+1$ colors

- Algorithm that has elements of both greedy algorithms and local search
 - Find an uncolored edge (u,v)
 - If possible, color it with one of the $\Delta+1$ colors (greedy)
 - Else, locally change some edge colors such that it is possible to color (u, v) (local search)
- Will not look into the details of the algorithm (page 49)

Theorem

- Theorem 2.23. There is a polynomial-time algorithm to find a $(\Delta + 1)$ -edge-coloring of a graph.